

# Victor Karani

[victor.karani@gmail.com](mailto:victor.karani@gmail.com)

## Spoiler Shield Mini Project

### Table of Contents

[Table of Contents](#)

[Summary and overall Approach](#)

[EDA](#)

[Data Overview](#)

[IMDB\\_reviews](#)

[IMDB\\_movie\\_details](#)

[Overall data management plan](#)

[Dataset size](#)

[Dataset size and Imbalance](#)

[Word Count](#)

[Word Cloud](#)

[TF-IDF -Cosine similarity](#)

[Model training and hyperparameter tuning](#)

[Model 1](#)

[Imbalanced data test \( not presented in the notebook\)](#)

[Model 1 Block diagram](#)

[Model 1 results](#)

[Model 2](#)

[Balanced sample distribution before training](#)

[Model 2 Block diagram](#)

[Hyperparameter tuning and interpretation with Model 2](#)

[Model 2: Sample size 1000, 6 epochs, lr 2e-5](#)

[Results for Model 2: Sample size 1000, 6 epochs, lr 2e-5](#)

[Classification report](#)

[Confusion Matrix Training vs validation](#)

[Training Vs Validation Loss](#)

[Model 2: increased learning rate \(2e-4\), 12 epochs, sample size of 5000](#)

[Classification report](#)

[Confusion Matrix](#)

[Training and validation loss](#)

[Model 2: increased learning rate \(2e-3\)](#)

[Conclusion: increased learning rate \(2e-3\)](#)

[Model 2: increased learning rate \(5e-4\), 12 epochs, sample size of 5000](#)

[Conclusion: increased learning rate \(5e-4\)](#)

[Model 2: bert-large-uncased](#)

[Conclusion: bert-large-uncased](#)

[Model 3](#)

[Model 3 Block diagram](#)

[Model 3: two tower with 10000 sample size](#)

[Conclusion:](#)

[Model 4](#)

[Model 4 Block diagram](#)

[Model 4 . Concatenate synopsis and summary](#)

[Conclusion: Model 4. Concatenate synopsis and Summary](#)

[Model 4 . Concatenate synopsis and summary, 50000 samples](#)

[Conclusion: Model 4. Concatenate synopsis and Summary](#)

[Model 4 Post training](#)

[Future Work](#)

## Summary and overall Approach

I chose the spoiler shield for my mini project. I went through several approaches to end up with the final result for this project.

## EDA

### Data Overview

The data for this project comes in 2 files:

- IMDB\_reviews.json
- IMDB\_movie\_details.json

## IMDB\_reviews

The **is\_spoiler** and **review\_text** are the important columns here, with **is\_spoiler** containing the **labels** and **review\_text** containing the actual reviews that are to be classified. We also need **movie\_id** as it will be used to join with the next data file.

## IMDB\_movie\_details

The **plot\_summary**, **plot\_synopsis** and **movie\_id** are the important columns here. The **plot\_summary** contains a movie description provided by the studio themselves, ideally written to be enticing but NOT to have spoilers. The **plot\_synopsis** contains movie details and can be assumed to have spoilers in it.

## Overall data management plan

I read the 2 files into pandas dataframes and then merged them on the **movie\_id** column.

## Dataset size

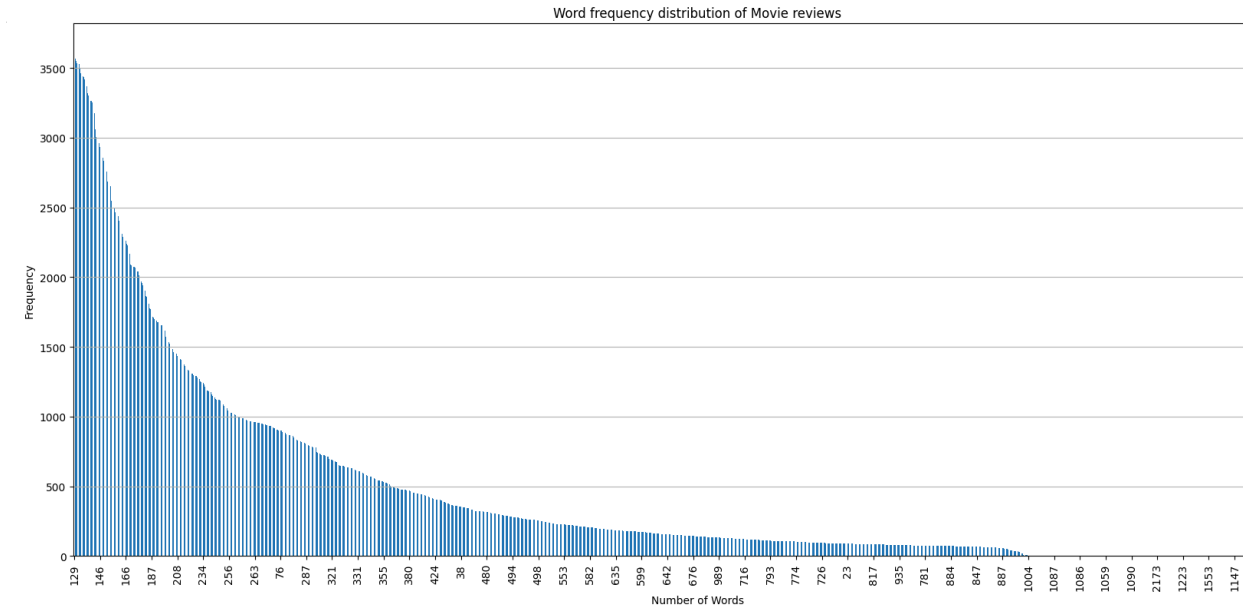
This is a very large dataset with 573906 rows and 13 columns after the merge.

## Dataset size and Imbalance

There are 422601 cases where **is\_spoiler** is **FALSE** and 150862 cases where **is\_spoiler** is **TRUE**. I will solve the giant size problem by only training on a portion of the dataset. I will solve the imbalance problem by creating a balanced dataset before training. This dataset is large enough that I don't believe I need synthetic data of any kind.

## Word Count

I performed a word count analysis. Looks like there is a majority (~ 3500 examples) of movie reviews at 129 words.



## Word Cloud

The word cloud gives a feeling of the most commonly used words. Not super useful, but interesting to look at:



## TF-IDF -Cosine similarity

I got a cosine similarity between the review\_text and plot\_summary of 0.39, which is pretty close.

## Model training and hyperparameter tuning

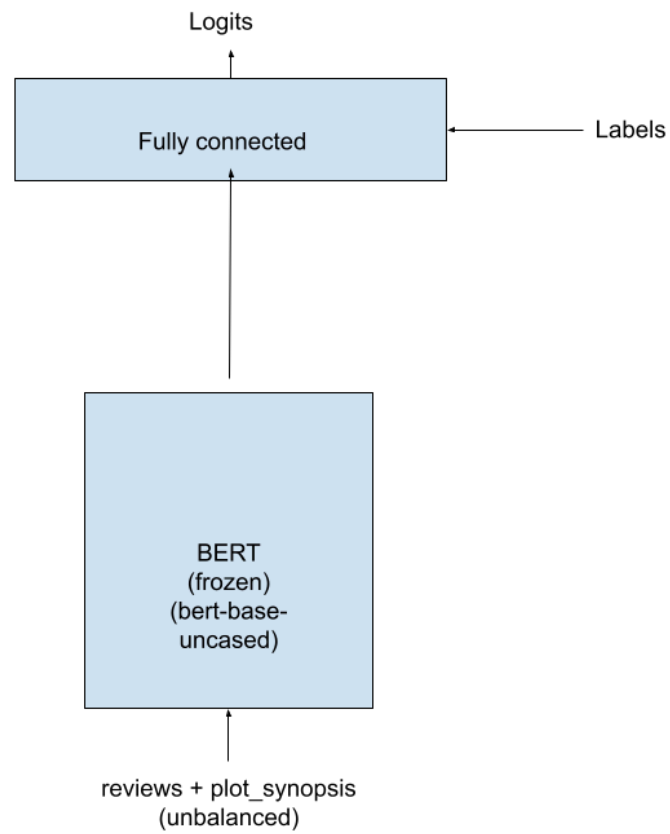
I had 3 distinct iterations for the approach to solve this problem. I also tried some experiments with unbalanced datasets

### Model 1

Imbalanced data test ( not presented in the notebook)

As an initial experiment, I tried at the very beginning to use unbalanced data.

## Model 1 Block diagram



## Model 1 results

The results were *spectacularly* bad. Below is a screenshot of what typically happened:

Epoch 1/4					
→	Validation Accuracy: 0.7478				
		precision	recall	f1-score	support
	0	0.75	1.00	0.86	2992
	1	0.47	0.01	0.01	1008
	accuracy			0.75	4000
	macro avg	0.61	0.50	0.43	4000
	weighted avg	0.68	0.75	0.64	4000
	Epoch 2/4				
	Validation Accuracy: 0.7480				
		precision	recall	f1-score	support
	0	0.75	1.00	0.86	2992
	1	0.50	0.01	0.02	1008
	accuracy			0.75	4000
	macro avg	0.62	0.50	0.44	4000
	weighted avg	0.69	0.75	0.64	4000
Epoch 3/4					
Validation Accuracy: 0.7488					
		precision	recall	f1-score	support
	0	0.75	1.00	0.86	2992
	1	0.62	0.01	0.02	1008
	accuracy			0.75	4000
	macro avg	0.68	0.50	0.44	4000
	weighted avg	0.72	0.75	0.64	4000
Epoch 4/4					
Validation Accuracy: 0.7485					
		precision	recall	f1-score	support
	0	0.75	1.00	0.86	2992
	1	0.58	0.01	0.01	1008
	accuracy			0.75	4000
	macro avg	0.67	0.50	0.43	4000
	weighted avg	0.71	0.75	0.64	4000

The accuracy seems high at 75% but, just look how uneven the precision, recall, and f1-score are between is\_spoiler 0 and 1 . This is so bad I didn't even bother with the confusion matrix here.

## Model 2

(not presented in the notebook)

I clearly need to balance the model.

It turns out that I don't have that much processing power even in Kaggle. I could only approach 10000 samples total. As a reminder, there are 422601 cases where is\_spoiler is **FALSE** and 150862 where is\_spoiler is **TRUE**.

So all I had to do was decide the maximum number of samples that I wanted to test, and then, pick an equal number of FALSE and TRUE is\_spoiler samples.

Below is a typical False/True (0/1) sample distribution prior to training

### Balanced sample distribution before training

This is how I balance the classes in code...

```
sample_size_per_class = 1000 # Number of samples per class
```

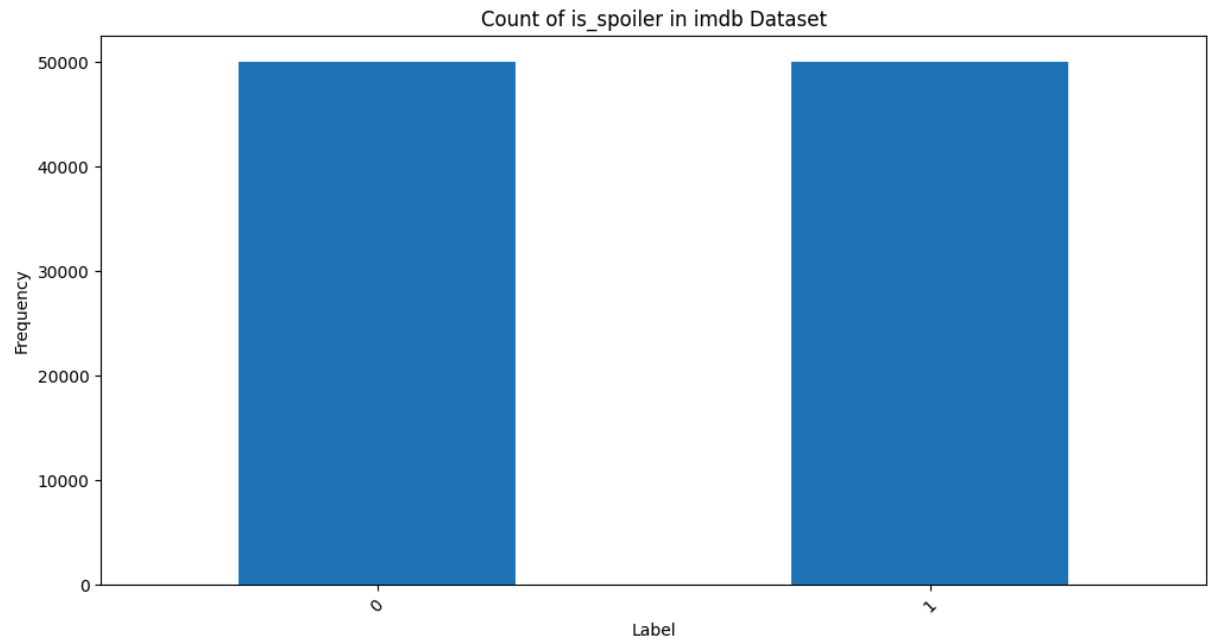
And

```
balanced_sample = (data.groupby('is_spoiler').apply(lambda x: x.sample(min(len(x), sample_size_per_class), random_state=42)).reset_index(drop=True))
```

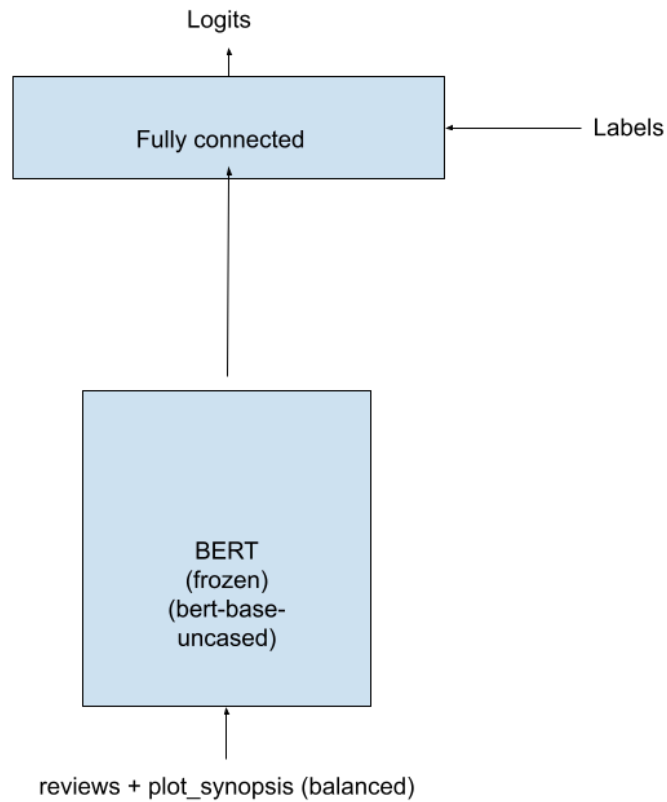
From here on in the document, I will refer to sample\_size\_per\_class as 'sample size'.

Below is one example where the sample size is 50000. As you can see is\_spoiler T and F are well balanced.





## Model 2 Block diagram



## Hyperparameter tuning and interpretation with Model 2

I did most of my hyperparameter tuning with Model 2, so I will go through all the hyperparameter tuning process here

Model 2: Sample size 1000, 6 epochs, lr  $2e-5$

Huggingface Model = **bert-base-uncased**

Sample size **1000**

Balanced (True vs false are equal in number)

num\_epochs=**6**

Learning\_rate  **$2e-5$**

## Results for Model 2: Sample size 1000, 6 epochs, lr 2e-5

### Classification report

```
[75]: # Final classification report  
print("Final Classification Report:")  
print(report)
```

Final Classification Report:


	precision	recall	f1-score	support
0	0.52	0.61	0.56	199
1	0.54	0.45	0.49	201
accuracy			0.53	400
macro avg	0.53	0.53	0.53	400
weighted avg	0.53	0.53	0.53	400

Classification report result:

Results (precision, recall, f1-summary are much better balanced than the previous unbalanced version. Accuracy is still pretty low, so we have a ways to go. Note that here, I am using a TINY sample size of 1000


## Confusion Matrix Training vs validation

```
[76]: # Display the confusion matrix
disp = ConfusionMatrixDisplay(confusion_matrix=cm)
disp.plot(cmap=plt.cm.Blues)
plt.title('Confusion Matrix for the DL Spoiler shield')
plt.xlabel('Predicted Label')
plt.ylabel('True Label')
plt.show()
```



## Training Vs Validation Loss

```
plt.show()
```



Validation seems to be bottoming out , so perhaps running this for more epochs might help. The tiny sample size of 1000 probably contributes to the very uneven validation loss seen here.

Model 2: increased learning rate ( $2e-4$ ), 12 epochs, sample size of 5000

Note: I first tried increasing the number of epochs (12) separately and found the performance was better. I also increased the sample size (5000) separately and found the performance also improved . Then, here, I also increased the learning rate.

**Huggingface Model = bert-base-uncased**

**Sample size 5000**

**Balanced (True vs false are equal in number)**

**num\_epochs=12**

**Learning\_rate  $2e-4$**

Classification report

```
# Final classification report
print("Final Classification Report:")
print(report)
```

Final Classification Report:


	precision	recall	f1-score	support
0	0.61	0.59	0.60	1012
1	0.60	0.61	0.60	988
accuracy			0.60	2000
macro avg	0.60	0.60	0.60	2000
weighted avg	0.60	0.60	0.60	2000

Classification report result:

Results (precision, recall, f1-summary and accuracy) are pretty good. Note that I am using a TINY sample size of 5000

## Confusion Matrix


```
[97]: # Display the confusion matrix  
disp = ConfusionMatrixDisplay(confusion_matrix=cm)  
disp.plot(cmap=plt.cm.Blues)  
plt.title('Confusion Matrix for the DL Spoiler shield')  
plt.xlabel('Predicted Label')  
plt.ylabel('True Label')  
plt.show()
```



### Confusion Matrix summary

The false positives and false negatives are higher than I would like. True positives and True negatives are at least pretty solid at ~ 60 %.

```
plt.show()
```



### Training and validation loss

The changes here (increased learning rate ( $2e-4$ ), 12 epochs, sample size of 5000) provide a decent loss reduction (bottoming out at 11 epochs)

### Model 2: increased learning rate ( $2e-3$ )

I reran the exact same test as before but tried to increase the learning rate even more.

**Huggingface Model = bert-base-uncased**

**Sample size 5000**

**Balanced (True vs false are equal in number)**

**num\_epochs=12**



**Learning\_rate  $2e-3$**



```
[129]: # Final classification report  
print("Final Classification Report:")  
print(report)
```

Final Classification Report:

	precision	recall	f1-score	support
0	0.62	0.57	0.59	1012
1	0.59	0.64	0.61	988
accuracy			0.60	2000
macro avg	0.60	0.60	0.60	2000
weighted avg	0.60	0.60	0.60	2000





Conclusion: increased learning rate ( $2e-3$ )

The classification report and Confusion matrix are similar. Unfortunately, I was too greedy. The validation loss is very unstable. I will try again with a lr slightly smaller than  $2e-3$ .

Model 2: increased learning rate ( $5e-4$ ), 12 epochs, sample size of 5000

**Huggingface Model = bert-base-uncased**

**Sample size 5000**

**Balanced (True vs false are equal in number)**

**num\_epochs=12**


**Learning\_rate  $5e-4$**


```
[195]: # Final classification report
print("Final Classification Report:")
print(report)
```

Final Classification Report:

	precision	recall	f1-score	support
0	0.61	0.59	0.60	1012
1	0.59	0.61	0.60	988
accuracy			0.60	2000
macro avg	0.60	0.60	0.60	2000
weighted avg	0.60	0.60	0.60	2000

```
[196]: # Display the confusion matrix
disp = ConfusionMatrixDisplay(confusion_matrix=cm)
disp.plot(cmap=plt.cm.Blues)
plt.title('Confusion Matrix for the DL Spoiler shield')
plt.xlabel('Predicted Label')
plt.ylabel('True Label')
plt.show()
```





Conclusion: increased learning rate ( $5e-4$ )

The classification report and Confusion matrix are similar. Validation loss is a little less smooth than  $2e-3$  but I think offers more chance to converge faster. I'll keep  $5e-4$  for future runs.

## Model 2: bert-large-uncased

I tried to use the bigger **bert-large-uncased** model to see if it would be dramatically better than the previously used bert-base-uncased.

**Huggingface Model = bert-large-uncased**

**Sample size 5000**

**Balanced (True vs false are equal in number)**


**num\_epochs=12**

**Learning\_rate  $5e-4$**


```
: # Final classification report
print("Final Classification Report:")
print(report)
```

Final Classification Report:

	precision	recall	f1-score	support
0	0.61	0.60	0.61	1012
1	0.60	0.61	0.60	988
accuracy			0.60	2000
macro avg	0.60	0.60	0.60	2000
weighted avg	0.60	0.60	0.60	2000







Conclusion: bert-large-uncased

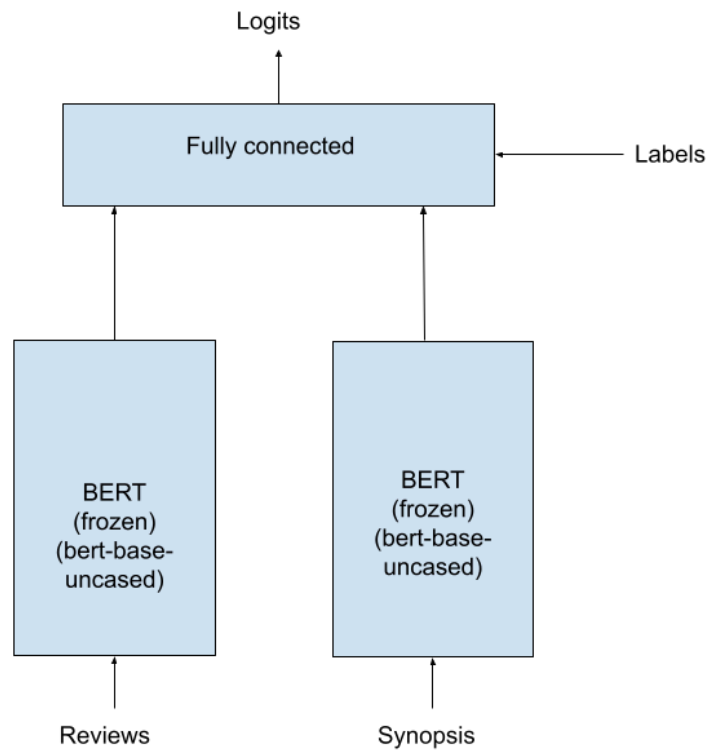
The classification report and Confusion matrix are similar. Validation loss is a little less smooth than  $2e-3$ . Overall, I don't think that the bert-large-uncased gives me dramatic improvements over bert-base-uncased. So I will keep bert-based-uncased.

## Model 3

(not presented in the notebook)

I think there are still more ways to improve the model. Next I tried the 2 tower model

Model 3 Block diagram



Model 3: two tower with 10000 sample size

Huggingface Model = bert-base-uncased

Sample size **100000**

Balanced (True vs false are equal in number)

num\_epochs=**4**

Learning\_rate 5e-4

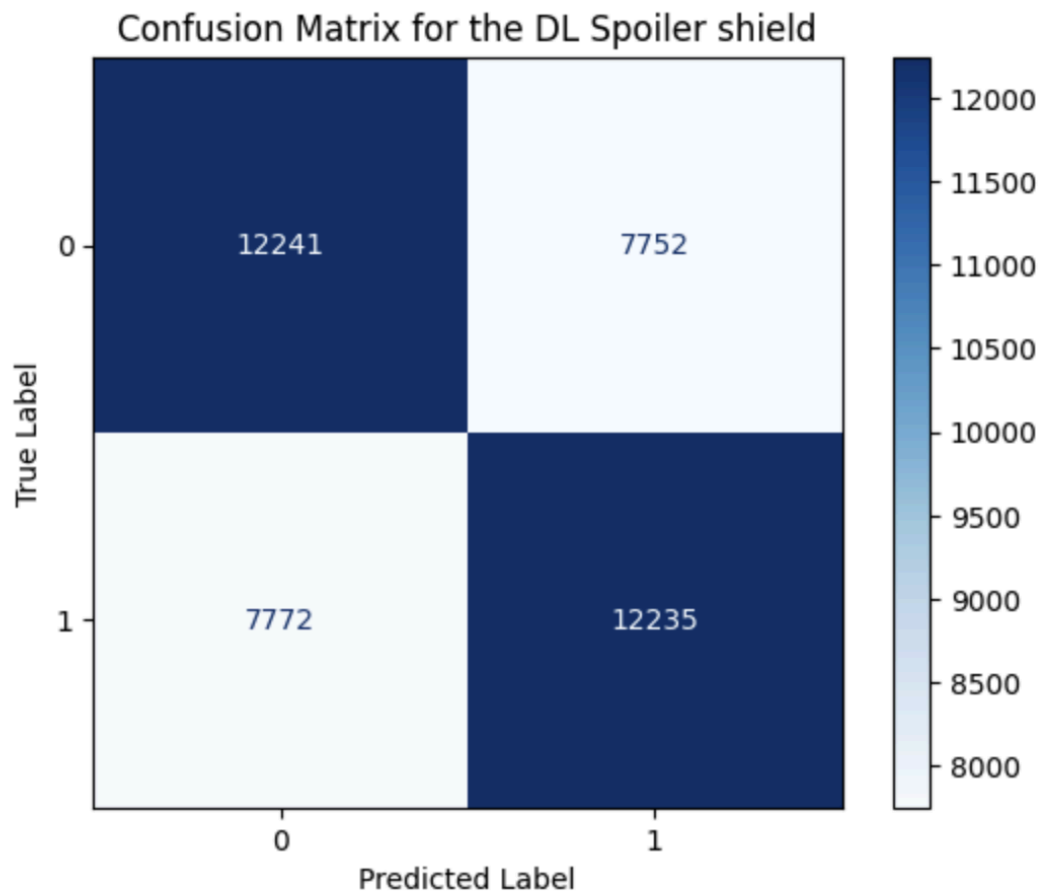
I used only 4 epochs for GPU time reasons.

```
[102]: # Final classification report  
print("Final Classification Report:")  
print(report)
```

Final Classification Report:


	precision	recall	f1-score	support
0	0.61	0.61	0.61	19993
1	0.61	0.61	0.61	20007
accuracy			0.61	40000
macro avg	0.61	0.61	0.61	40000
weighted avg	0.61	0.61	0.61	40000

```
[103]: # Display the confusion matrix
disp = ConfusionMatrixDisplay(confusion_matrix=cm)
disp.plot(cmap=plt.cm.Blues)
plt.title('Confusion Matrix for the DL Spoiler shield')
plt.xlabel('Predicted Label')
plt.ylabel('True Label')
plt.show()
```



```
[104]: cm_normalized = cm.astype('float') / cm.sum(axis=1)[:, np.newaxis]
disp_normalized = ConfusionMatrixDisplay(confusion_matrix=cm_normalized)
disp_normalized.plot(cmap=plt.cm.Blues)
plt.title('Normalized Confusion Matrix for the DL Spoiler shield')
plt.xlabel('Predicted Label')
plt.ylabel('True Label')
plt.show()
```

Normalized Confusion Matrix for the DL Spoiler shield





## Conclusion:

The two tower accuracy is at 61% which is just about the top level that I was able to get with this project.

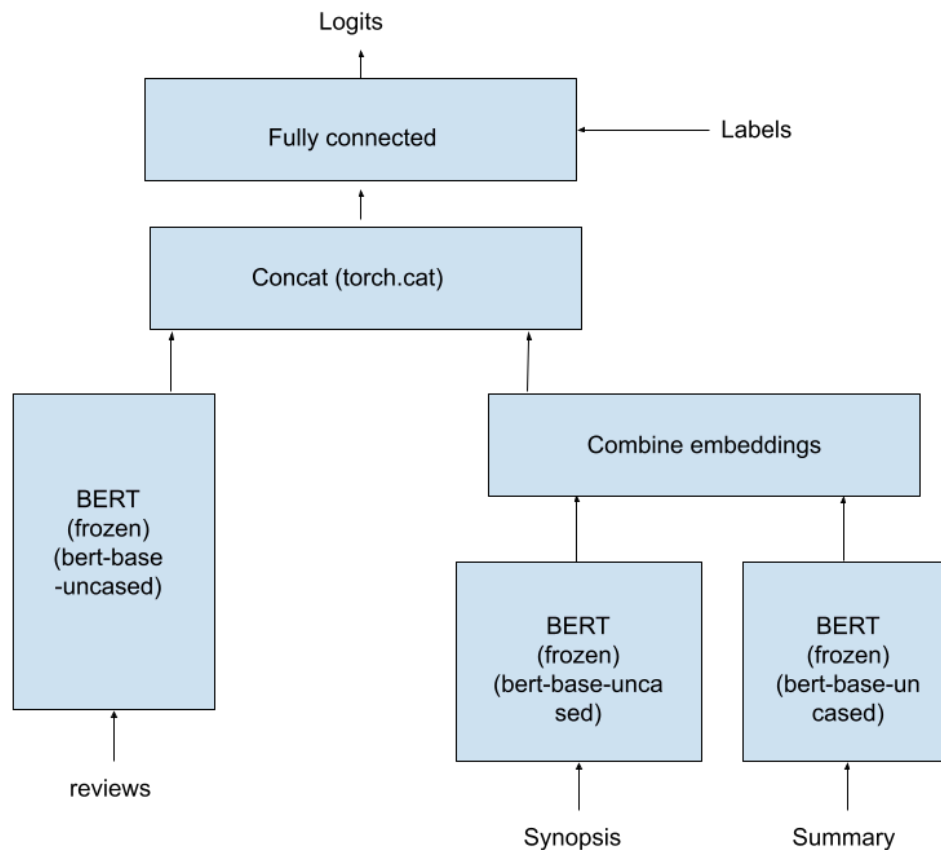
The large sample size and 2 tower method enables the model to converge to validation loss below .65 but still slightly underfit. Unfortunately, I didn't have enough GPU time to go over 4 epochs.

## Model 4

(presented in this notebook)

I am aware that the 3 BERT heads described below involve more overhead than the classic 2 tower model, but I wanted to investigate it and see if I could get better performance.

## Model 4 Block diagram



I tried several methods for combining the synopsis and summary

- Subtracting the synopsis and summary embeddings
- Adding the synopsis and summary embeddings
- Concatenating the synopsis and summary embeddings

None of these seemed to be dramatically better than the 2 tower method but the concatenation method seemed to provide the best results, so I went with it.

## Model 4 . Concatenate synopsis and summary

Huggingface Model = bert-base-uncased

Sample size **1000**

Balanced (True vs false are equal in number)  
num\_epochs=4  
Learning\_rate 5e-4


```
✓ [79] # Final classification report  
0s print("Final Classification Report:")  
    print(report)
```

```
➞ Final Classification Report:
```

	precision	recall	f1-score	support
0	0.51	0.59	0.55	199
1	0.52	0.44	0.48	201
accuracy			0.52	400
macro avg	0.52	0.52	0.51	400
weighted avg	0.52	0.52	0.51	400


---








Normalized Confusion Matrix for the DL Spoiler shield





Conclusion: Model 4. Concatenate synopsis and Summary

I limited this model to 1000 sample size because I just wanted a quick test without using too much CPU time. The accuracy is actually not that great. It is not really better than the 2 tower model.

Model 4 . Concatenate synopsis and summary, 50000 samples

This is the model presented in the document

Huggingface Model = bert-base-uncased

Sample size **50000**

Balanced (True vs false are equal in number)

num\_epochs=4

Learning\_rate 5e-4


```
[39]: # Final classification report
print("Final Classification Report:")
print(report)
```

```
Final Classification Report:
              precision    recall  f1-score   support

     0       0.63         0.60         0.61      10035
     1       0.61         0.64         0.63       9965


 accuracy                   0.62      20000
 macro avg                  0.62         0.62         0.62      20000
 weighted avg              0.62         0.62         0.62      20000
```


```
[40]: # Display the confusion matrix
disp = ConfusionMatrixDisplay(confusion_matrix=cm)
disp.plot(cmap=plt.cm.Blues)
plt.title('Confusion Matrix for the DL Spoiler shield')
plt.xlabel('Predicted Label')
plt.ylabel('True Label')
plt.show()
```



```
[41]: cm_normalized = cm.astype('float') / cm.sum(axis=1)[:, np.newaxis]
disp_normalized = ConfusionMatrixDisplay(confusion_matrix=cm_normalized)
disp_normalized.plot(cmap=plt.cm.Blues)
plt.title('Normalized Confusion Matrix for the DL Spoiler shield')
plt.xlabel('Predicted Label')
plt.ylabel('True Label')
plt.show()
```

Normalized Confusion Matrix for the DL Spoiler shield





## Session options

ACCELERATOR

GPU T4 x2

Quota: 24:14 / 30 hrs

LANGUAGE

One training run for 4 epochs used up all this kaggle compute time!!

Conclusion: Model 4. Concatenate synopsis and Summary

The 3 tower model is very expensive in CPU time but not demonstrably much better than the 2 tower model which I did earlier. However, I think If I had more time, I could extract much more performance from this model.

## Model 4 Post training

I also saved this model and used it on brand new data and found that the Post training performance was consistent with the training

```
: # Final Post training classification report
print("Final Post training Classification Report:")
print(report)
```

Final Post training Classification Report:

	precision	recall	f1-score	support
0	0.62	0.61	0.61	1000
1	0.61	0.63	0.62	1000
accuracy			0.62	2000
macro avg	0.62	0.62	0.62	2000
weighted avg	0.62	0.62	0.62	2000




```

|: cm_normalized = cm.astype('float') / cm.sum(axis=1)[:, np.newaxis]
   disp_normalized = ConfusionMatrixDisplay(confusion_matrix=cm_normalized)
   disp_normalized.plot(cmap=plt.cm.Blues)
   plt.title('Normalized Confusion Matrix for the DL Spoiler shield')
   plt.xlabel('Predicted Label')
   plt.ylabel('True Label')
   plt.show()

```

Normalized Confusion Matrix for the DL Spoiler shield



Future Work

I would love to have had more time to pursue the 3- tower model. I think I could have extracted more performance there.

In retrospect, I should probably have used a lighter model like ALBERT to do initial training then switched to bert proper. This could have probably saved me some time and CPU.

I also tried BERT for NSP but many of the texts kept getting truncated, invalidating the runs. I am sure there is a way to fix this but didn't have time to pursue this.

The is\_spoiler data is probably noisy. I think that overall the system performance is good.

Overall, this was a great project. It got me hands on with BERT. I can't wait to use more sophisticated models in upcoming projects.